
Identification of Latent Variables From Their Footprint In Bayesian Network Residuals

Anonymous Author(s)

Affiliation

Address

email

Abstract

Graph-based causal discovery methods aim to capture conditional independencies consistent with the observed data, differentiating causal correlations from indirect or induced ones. Successful construction of graphical models of data depends, among others, on the assumption of causal sufficiency. For partially observed data, graphical model structures may become arbitrarily incorrect, and effects implied by such models may be wrongly attributed, carry wrong magnitude, or mis-represent direction of correlation. Wide applicability and application of graphical models to increasingly less and less curated "big data" highlights the need for continued attention to the unobserved confounder problem.

We present a novel method that aims to control for the latent structure of the data by iteratively deriving proxies for the latent space from the residuals of the inferred graphical model. Under mild assumptions, our method improves structural inference. In addition, when the model is being used to predict outcomes, this method un-confounds the coefficients on the parents of the outcomes and leads to improved predictive performance when out-of-sample regime is very different from the training data. We show that such improvement of the predictive model is intrinsically capped and cannot be improved beyond a certain limit as compared to the confounded model. Furthermore, we propose an algorithm for computing a ceiling for the dimensionality of the latent space which may be useful in future approaches to the problem.

1 Introduction

Construction of graphical models (GMs) at its heart pursues two related objectives: accurate inference of the structure of conditional independencies and construction of predictive models for outcomes of interest with the purpose of estimating average causal effect (ACE) with respect to interventions (Pearl [2000], Hernán and Robins [2006]). These two distinct goals mandate that certain identifiability conditions for both types of tasks be met. Of particular interest to this work is the condition of *causal sufficiency* (Spirtes et al. [1993]) - namely, whether all of the relevant predictors have been observed. When this condition is not met, both accurate GMs and correct ACEs are hard to infer.

Hitherto, the literature has largely focused on addressing the subset of problems when ACE could be estimated reliably in the absence of this guarantee (such as when conditional exchangeability holds - e.g., Hernán and Robins [2006]).

We aim to show that there exist circumstances when causal sufficiency can be asymptotically achieved, and thus exchangeability ensured, even when the causal drivers of outcome are confounded by a number of latent variables. This can be achieved when the confounding is **pleiotropic** - when the latent variable affects a "large enough" number of variables, some driving an outcome of interest and others not (Anandkumar et al. [2013]). Notably, this objective cannot be achieved when confounding affects only the variables of interest and their causal parents (D'Amour [2019]).

Submitted to 33rd Conference on Neural Information Processing Systems (NeurIPS 2019). Do not distribute.

Intuitively, presence of broad unobserved confounding gives rise to violation of conditional independence among the affected variables downstream from the latent confounder. Likelihood methods for GM construction aim to minimize unexplained variance for all variables in the network by accounting for conditional independencies in the data (Pearl [2000], Friedman and Koller [2013]). Lack of observability of a causally important variable will induce dependencies among its descendants in the graph that cannot be fully ascribed to any single "heir" of the latent variable except by chance due to noise. Such unexplained interdependency results in model residuals correlated with the latent variable and not fully explained by any putative graph parents, and thus to inferred connectivity that's "excessive" as compared to the true network and in appearance of (near-)cliques (Elidan et al. [2001]).

Previously, methods have been proposed for inferring latent variables affecting GMs by the means of EM (as far back as Friedman and others [1997], Friedman [1998]). However, for a large enough network local gradients do not provide a reliable guide, nor do they address the cardinality of the latent space. Methods for using near-cliques for detection of latent variables in directed acyclic graphs (DAGs) (Elidan et al. [2001]), including with gaussian graphical models (GGMs), have been proposed (Silva et al. [2006]) that address both problems by analyzing near-cliques in DAGs. A method related to ours has been proposed for calculating latent variables in a greedy fashion in linear and "invertible continuous" networks, and relating such estimates to observed data to speed up structure search and enable discovery of hidden variables (Elidan et al. [2007]).

2 Background And Notation

We are going to concern ourselves with a factorized joint probability distribution over a set of observed and latent variables (see Table 1 for notation).

Set	Meaning	Indexing
S	samples	$S_i, i \in \{1, \dots, s\}$
V	observed predictor variables	$V_j, j \in \{1, \dots, v\}$
U	unobserved predictor variables	$U_l, l \in \{1, \dots, u\}$
O	outcomes (sinks)	$O_k, k \in \{1, \dots, o\}$
D	$\{V, O\}$ - observable data	
D_u	$\{V, O, U\}$ - implied data	
θ	parameters	$\theta_i, i \in \{1, \dots, t\}$
P^N	parents of variable N	$P_i^N, i \in \{1, \dots, p\}$
C^N	children of variable N	$C_q^N, q \in \{1, \dots, c\}$
G	graph over D	
G_u	graph over D_u	

Table 1: Notation

Assume that the joint distribution D (D_u) is factorized as a directed acyclic graph, G (G_u). We will consider individual conditional probabilities describing nodes and their parents, $P(V|parents(V), \theta)$, where θ refers to the parameters linking the parents of V to V . $\hat{\theta}$ will refer to an estimate of these parameters. We will furthermore assume that G is constructed subject to regularization and using unbiased estimators for $P(V|parents(V), \hat{\theta})$. We will further assume that D_u plus any given constraints are sufficient to infer the true graph up to markov equivalence. For convenience, we'll focus on the actual true graph's parameters, so that, using unbiased estimators, $E[\hat{\theta}_m|D_u] = \theta_m, \forall m$.

Mirroring D (or D_u), we will define a matrix R (or R_u) of the same dimensions - $s \times (v + o)$ (or $s \times (v + o + u)$) - that captures the residuals of modeling every variable $N \in \{V, O, (U)\}$ via G (or G_u). In the linear case, these would be regular linear model residuals, but more generally we will consider probability scale residuals (PSR, Shepherd et al. [2016]). That is, we define $R[i, j] = PSR(P(V_j|parents(V_j), \hat{\theta}_j)|D[i, j])$, the residuals of V_j given its graph parents. Notice that the use of probability-scale residuals allows us to define R and R_u for all ordinal variable types, up to rank-equivalence.

3 Algorithm

3.1 Gaussian Graphical Models (GGMs)

Recall that, for some $V_j \in \{V, U, O\}$, $P_k^{V_j}$ denotes the k th parent of V_j . For GGMs, we can write down a fragment of any DAG G as a linear equation:

$$V_j = \beta_{j0} + \beta_{j1}P_1^{V_j} + \dots + \beta_{jp}P_p^{V_j} + \xi_j, \quad \xi_j \sim \mathcal{N}(0, \sigma_j). \quad (1)$$




Figure 1: Graph G_u . U influences the outcomes O , and a number of predictors V , confounding many of the $V_j \rightarrow O_k$ relationships. Gray nodes are affected by U .




Figure 2: Graph G . With U latent, the graph adjusts, introducing spurious edges.

For example, consider O_1 in Figure 1. We can write:

$$O_1 = \beta_0 + \beta_6 V_6 + \beta_5 V_5 + \beta_4 V_4 + \beta_7 V_7 + \beta_u U + \xi_1, \quad \xi_1 \sim \mathcal{N}(0, \sigma_{O_1}). \quad (2)$$

For any variable N that has parents in G_u , we can group variables in P^N into three subsets: $X_U \in \{P^U, C^U\}$, $X_{\mathcal{H}} \notin \{P^U, C^U\}$, and the set U itself, and write down the following general form using matrix notation:

$$N = \beta_{N0} + B_U X_U + B_{\mathcal{H}} X_{\mathcal{H}} + \beta_U U + \xi_N, \quad \xi_N \sim \mathcal{N}(0, \sigma_N). \quad (3)$$

Explicit dependence of N on U happens when $\beta_U \neq 0$.

Now consider G - the graph built over the variables $\{V, O\}$ excluding the latent space U . Note that if we deleted U and its edges from G_u without rebuilding the graph, Equation 3 from G_u would read:

$$N = \beta_{N0} + B_U X_U + B_{\mathcal{H}} X_{\mathcal{H}} + R_N + \xi_N. \quad (4)$$

The residual term R_N is simply equal to the direct contribution of U to N . The network G would have to adjust to the missingness of U (e.g., Figure 2 vs Figure 1). As a result, R_N will be partially substituted by other variables in $\{P^U, C^U\}$. Still, unless U is completely explained by $\{P^U, C^U\}$ (as described in D'Amour [2019]) and in the absence of regularization (when a high enough number of covariates may lead to such collinearity), R_N will not fully disappear in G . Hence, even after partially explaining the contribution of U to N by some of the parents of N in G ,

$$R_N = \beta_0 + \beta_1 U + \xi_N. \quad (5)$$

91

Samples	Variables				Samples	Residuals			
	V_{11}	V_{12}	\dots	V_{1v}		R_{11}	R_{12}	\dots	R_{1v}
	\vdots	\vdots	\vdots	\vdots		\vdots	\vdots	\vdots	\vdots
	V_{s1}	V_{s2}	\dots	V_{sv}		R_{s1}	R_{s2}	\dots	R_{sv}

Table 2: The training data frame (left) implies a matching residual data frame (right) once the joint distribution of all variables is specified via a graph and its parameterization

92

Therefore, the columns in the residuals table corresponding to G (Table 2) that represent the parents and children of U will contain residuals collinear with U :

$$\begin{aligned} R_1 &= \beta_{10} + \beta_{11} U + \xi_1 \\ &\vdots \\ R_k &= \beta_{k0} + \beta_{k1} U + \xi_k. \end{aligned} \quad (6)$$

Rearranging and combining,

$$U = \beta_i^* R_1 + \dots + \beta_k^* R_k + \xi = BR + \xi. \quad (7)$$

Equation 7 tells us that, for graphical gaussian models, components of U are obtainable from linear combinations of residuals, or principal components (PCs) of the residual table R . In other words,

98 U is identifiable by principal component analysis (PCA). Whether the residuals needed for this
 99 identification exist depends on the *expansion property* as defined in Anandkumar et al. [2013].

100 Note that this approach is similar, in the linear case, to that in Elidan et al. [2007] except insofar as
 101 we show the global principal components of the residual matrix to be optimal for the discovery of the
 102 whole latent space of a given DAG assuming its structure is already correct, and we therefore couple
 103 structure inference and latent space discovery via EM as separate rather than interleaved steps (1).

104 4 Confounding of outcomes

105 Aside from the inference of the **exact** network - an interesting exercise that is probably futile in
 106 any practical application owing to the complexity of the inference problem - the most important
 107 utility of causal modeling is to propose suitable predictors, as well as predictors of predictors, for
 108 outcomes of interest. These "predictors of predictors" may have practical importance - for instance,
 109 when developing a drug, direct predictors of an outcome, say V_6 and O_1 from Figure 1), may not be
 110 druggable, but some of the mediators of treatment upstream of direct predictors, such as V_5 , may turn
 111 out to be promising drug targets. Moreover, in the presence of latent confounding, coefficients of
 112 predictors of O_1 may be indeterminate (e.g. Hernán and Robins [2006]). For example, U induces
 113 correlation among V_3 , V_4 , V_5 , and V_6 even when these variables are conditionally independent given
 114 U .

The extent of this latter problem can be quantified. Suppose we model O_1 without controlling for U
 (2):

$$O_1 = \beta_0 + \beta_3 V_3 + \beta_4 V_4 + \beta_5 V_5 + \beta_6 V_6 + \dots$$

Let's set the coefficient of determination for the model

$$V_3 = \alpha_0 + \alpha_4 V_4 + \alpha_5 V_5 + \alpha_6 V_6 + \dots$$

115 equal to ρ_3^2 . Then the estimated variance of β_3 in the presence of collinearity can be related to that
 116 when collinearity is absent via the following formula (Rawlings et al. [1998]):

$$var(\bar{\beta}_3) = var(\beta_3) \frac{1}{1 - \rho_3^2} \propto \frac{1}{1 - \rho_3^2}. \quad (8)$$

117 Formula 8 describes the *variance inflation factor* (VIF) of β_3 . Note that $\lim_{\rho \rightarrow 1} \frac{1}{1 - \rho^2} = \infty$, so even
 118 mild collinearity induced by latent variables can severely distort coefficient values and signs, and
 119 thus estimation of ACE. The method outlined above will reduce the VIFs of coefficients related to
 120 outcomes and thus make all *causal* statements relating to outcomes, such as calculation of ACE, more
 121 reliable, since by controlling for \bar{U} - the estimate of U - in the network,

$$\lim_{(U - \bar{U}) \rightarrow 0} var(\bar{\beta}_i) = var(\beta_i). \quad (9)$$

122 Can we hope to reach this limit? Consider an output O_j . While it is difficult to describe the limit
 123 of error on the coefficients of the predictors of O_j , it is straightforward to put a ceiling on the
 124 improvement in the likelihood obtainable from modeling U and approximating G_u with $G_{\bar{u}}$. Suppose
 125 we eventually model U as a linear combination of a set of variables $X \subset V$, and denote by $X \setminus W$
 126 the set difference: members of X not in W . Then for any outcome O_i predicted by a set of variables
 127 W in the graph G and in truth predicted by the set $Z + U$, we can contrast three expressions (from G
 128 and $G_{\bar{U}}$ respectively):

$$\begin{aligned} O_i &= \beta_{i0} + B_W W(a) + \xi_i & (a) \\ O_i &= \beta_{i0} + B_W W + B_{X \setminus W}(X \setminus W) + \xi_i & (b) \\ O_i^U &= \beta_{i0}^U + B_Z^U Z + B^U U + \xi_i & (c). \end{aligned} \quad (10)$$

129 Model (a) is the model that was actually accepted, subject to regularization, in G . Model (b) is
 130 the "oracular" model of O_i that controls for U non-parsimoniously by controlling for all variables
 131 affected by U and not originally in the model. The third model, (c), is the ideal parsimonious model
 132 when U is known. We can compare the quality of these models by Bayesian Score, and the full score,
 133 in large sample sizes, can be approximated by BIC - the Bayesian Information Criterion (Koller and
 134 Friedman [2009]). We assume that the third of these equations would have the lowest BIC (being the
 135 best model), and the first being the second highest, since we know that the set of variables $X \setminus W$

136 didn't make it into the first equation subject to regularization by BIC. Assuming n samples,

$$\begin{aligned} BIC(O_i = \beta_{i0} + B_W W + \xi_i) &= b_a & (a) \\ BIC(O_i = \beta_{i0} + B_W W + B_{X \setminus W}(X \setminus W) + \xi_i) &= b_b = b_c + |X \setminus W| \log(n) & (b) \\ BIC(O_i^U = \beta_{i0}^U + B_Z^U Z + B^U U + \xi_i) &= b_c & (c). \end{aligned} \quad (11)$$

137 The "oracular model" - model (b) - includes all of the true predictors of O_j . Therefore its score will
138 be the same as that of the true model - model (c) - plus the BIC penalty, $\log(n)$, for each extra term,
139 minus the cost of having U in the true model (that is, the cardinality of the relevant part of the latent
140 space). We know that the extra information carried by this model was not big enough to improve
141 upon model (a), that is $b_a < b_c + k \log(n)$ for some k . Rearranging:

$$b_c - b_a > -k \log(n). \quad (12)$$

142 Any improvement in $G_{\bar{U}}$ owing to modeling of \bar{U} cannot, therefore, exceed $k \log(n)$ logs, where
143 $k = |X \setminus W| - |U|$: the information contained in the "oracular" model is smaller than its cost.

144 Although the available improvement in predictive power is also capped in some way, it is still
145 important to aim for that limit. The reason is, correct inference of causality, especially in the presence
146 of latent variables, is the only way to ensure transportability of models in real-world (heterogeneous-
147 data) applications (see, e.g., Bareinboim and Pearl [2016]).

148 Up to here, our method is a generalization of work presented in Anandkumar et al. [2013], where the
149 authors show that under some assumptions the latent space can be learned exactly, and which is also
150 related to that described by Wang and Blei [2018]. However, we do not require that the observables
151 be conditionally independent given the latent space and instead *generate* such independence by the
152 use of causal network's residuals, which are, of course, conditionally independent of each other *given*
153 *the graph and the latent space*. However, since the network among the observables is undefined
154 in the beginning, the structure of the observable network must be learned at the same time as the
155 structure of the latent space, which leads us to the iterative/variational bayes approach presented in
156 Algorithm 1. Lastly, the use of the entirety of the residual space is different from the work described
157 in Elidan et al. [2007], where local residuals are pursued with the goal to accelerate structure learning
158 while simultaneously discovering the latent space.

159 4.1 Gaussian Graphical Models With Interactions

160 In the presence of interactions among variables in a GGM, equation 5 expressing the deviation
161 of residuals from Gaussian noise may acquire higher-order terms due to interactions among the
162 descendants of the latent space U :

$$R_N = \beta_0 + \beta_1 U + \beta_2 U^2 + \beta_3 U^3 + \dots \quad (13)$$

163 Assuming interactions up to k th power are present in the system being modeled, residuals for each
164 variable may have up to k terms in the model matrix described by equation 13, and if interactions
165 among variables in the latent space U also exist, the cardinality of the principal components of the
166 residuals may far exceed the cardinality of the underlying latent space. Nevertheless, it may be
167 possible to reconstruct a parsimonious basis vector by application of regularization and nonlinear
168 approaches to latent variable modeling, such as nonlinear PCA (e.g. using methods from Karatzoglou
169 et al. [2004]), or autoencoders (Louizos et al. [2017]) as will be discussed below.

170 4.2 Generalization to nonlinear functions

171 We can show that linear PCA will suffice for a set of transformations broader than GGMs without
172 interactions. In particular, we will focus on nonlinear but homeomorphic functions within the
173 Generalized Additive Model (GAM) family. When talking about multiple inputs, we will require
174 that the relationship of any output variable to any of the covariates in equation 5 is homeomorphic
175 (invertible), and that equation 7 can be marginalized with respect to any right-hand-side variable as
176 well as to the original left-hand side variable. For such class of transformations, mutual information
177 between variables, such as between a single confounder U and some downstream variable N , is
178 invariant (Kraskov et al. [2004]). Therefore, residuals of any variable N will be rank-correlated
179 to $\text{rank}(U)$ in a transformation-invariant way. Further, spearman rank-correlation, specifically, is
180 defined as pearson correlation of ranks, and pearson correlation is a special case of mutual information

181 for bivariate normal distribution. Therefore when talking about mutual information between ranks of
 182 arbitrarily distributed variables, we can use our results for the GGM case above.

183 Thus, equation 5 will apply here with some modifications:

$$\text{rank}(R_N) = \beta_0 + \beta_1 \text{rank}(U) + \xi_N. \quad (14)$$

184 Since a method has been published recently describing how to capture rank-equivalent residuals (aka
 185 probability-scale residuals, or PSR) for any ordinal variable (Shepherd et al. [2016]), we can modify
 186 the equation 7 to reconstruct latent space up to rank-equivalence when interactions are absent from
 187 the network.

$$\text{rank}(U) = \frac{1}{\beta_i} \text{rank}(R_i) + \frac{1}{\beta_j} \text{rank}(R_j) + \dots + \xi. \quad (15)$$

188 When U consists of multiple variables that are independent of each other, the relationship between N
 189 and U can be written down using the mutual information chain rule (MacKay [2003]) and simplified
 190 taking advantage of mutual independence of the latent sources:

$$I(N; U) = I(N; U_1, U_2, \dots, U_U) = \sum_{i=1}^u I(N; X_i | X_{i-1}, \dots, X_1) = \sum_{i=1}^u I(N; X_i). \quad (16)$$

191 If interactions among U are present, it may still be possible to approximate the latent space with a
 192 suitably regularized nonlinear basis, but we do not, at present, know of specific conditions when this
 193 may or may not work. Novel methods for encoding basis sets, such as nonlinear PCA (implemented
 194 in the accompanying code), autoencoders, and others, may be brought to bear to collapse the
 195 linearly independent basis down to non-linearly independent (i.e. in the mutual information sense)
 196 components.

197 While approximate inference of latent variables for GMs built over invertible functions had been
 198 noted in Elidan et al. [2007], the above method gives a direct rank-linear approach leveraging the
 199 recently-proposed PSRs.

200 4.3 Generalization to categorical variables

201 In principle, PSRs can be extended to the case of non-ordinal categorical variables by modeling binary
 202 in/out of class label, deviance being correct/false. These models would lack the smooth gradient
 203 allowed by ranks and would probably converge far worse and offer more local minima for EM to get
 204 stuck in.

205 5 Implementation

206 Algorithm 1 below describes our approach to learning the latent space and can be viewed as a type of
 207 an expectation-maximization algorithm, possibly nested, if EM is used to learn the DAG at each step.

208 How do we learn $\bar{U} = f(R_{\bar{U}})$? In the linear case, we can use PCA, as described above, and in the
 209 non-linear case, we can use non-linear PCA, autoencoders, or other methods, as alluded to above as
 210 well. However, the linear case provides a useful constraint on dimensionality, and this constraint can
 211 be derived quickly. A useful notion of the ceiling constraint on the linear latent space dimensionality
 212 can be found in Gavish and Donoho [2014]. From a practical standpoint, the dimensionality can be
 213 even tighter, and we propose a permutation-test-based method for inferring $\text{ceiling}(|U|)$ in Algorithm
 214 2.

215 The integral should converge faster than the count of times variance explained by U^* on true residuals
 216 exceeds that obtained from shuffled residuals, but the permutation test approach of PCA cardinality is
 217 also workable, albeit with more iterations. Note that it is necessary to correct for the number of tests
 218 performed, and that we use Bonferroni correction as a simple and conservative stand-in. Alternatively,
 219 networks built using structural priors of the form proposed in Friedman and Koller [2013] may not
 220 need to perform this step.

221 Our algorithm 2 is different from that proposed in Elidan and Friedman [2005]. If we consider
 222 $Y = f(X)$, where Y are all DAG outputs and X are all inputs, and f is a suitably parameterized

DAG, PCA over residuals (linear or not) normalized by PCA over shuffled residuals provides a measure of "compressibility" of residual space. In other words, we propose a minimum description length algorithm for detecting the latent variables so that the residual space is no longer compressible.

Algorithm 1: Learning \tilde{U} from structure residuals via EM

Data: The set of observed variables $\{V, O\}$
Result: Graph $G_{\tilde{U}}(V, O, \tilde{U})$
Construct $G = G(V, O)$
Compute $S_0 = BIC_G$
Estimate $\tilde{U} = f(R)$
Construct $G_{\tilde{U}} = G(V, O, \tilde{U})$
Compute $S_{\tilde{U}} = BIC(G_{\tilde{U}})$
while $S_{\tilde{U}} - S_0 > \epsilon$ **do**
 Set $S_0 = S_{\tilde{U}}$
 Calculate $R_{\tilde{U}}$
 Set \tilde{U} to arbitrary constant values
 foreach child node $C \in G_{\tilde{U}}, C \notin \tilde{U}$ **do**
 Set parents to training data
 $C = C|parents(C)$
 Set $R_C = PSR(C, C)$
 end
 Estimate $\tilde{U} = f(R_{\tilde{U}})$
 Construct $G_{\tilde{U}} = G(V, O, \tilde{U})$
 Compute $S_{\tilde{U}} = BIC(G_{\tilde{U}})$
end

Algorithm 2: Inferring linearly optimal \tilde{U} and assessing its cardinality by permutations

Data: The set of residuals R_G from modeling D with $G_{\tilde{U}}$
Result: Linear approximation to \tilde{U}
Set significance threshold α (e.g. $\alpha = 0.05$)
Learn $\tilde{U} = PCA(R_G)$
Calculate column-wise variance explained V_E^0 for \tilde{U}^*
Set $V_E^* = 0 \times rank(R_G)$, matrix of variances explained by shuffling
while $se(\tilde{U}) > \epsilon$ **do**
 $R_G^* = shuffle(R_G)$ (column-wise)
 Calculate $\tilde{U}^* = PCA(R_G^*)$
 Calculate V_E^* for \tilde{U}^*
 Concatenate row-wise: $V_E^* = \{V_E^*, V_E^*\}$
 Fit $B(i)$ beta distributions to each column i of V_E
 For each column i of \tilde{U} , calculate:

$$P(V_E^0(i)|V_E^*(i)) = \lim_{|V_E^0(i)| \rightarrow \infty} \frac{|V_E^0(i) > V_E^*(i)|}{|V_E^0(i)|} \approx 1 - \int_{-\infty}^{V_E^0(i)} PDF(B_i)$$

end
 $P(V_E^0(i)|V_E^*(i)) = P(V_E^0(i) \sim V_E^*(i)) \times rank(V_E)$
Drop $V_E^0(i)$ for which $P(V_E^0(i) \sim V_E^*(i)) > \alpha$

6 Numerical Demonstration

To illustrate the algorithms described in the previous sections we generated synthetic data from the network shown in Figure 3 where two variables V_1 and V_2 drive an outcome Z . Two confounders U_1 and U_2 affect both the drivers and the outcome, as well as many additional variables that do not affect the outcome Z . The coefficient values in the network were chosen so that faithfulness was achieved and that the structure and coefficients were approximately recovered when all variables were observed.

The underlying network inference needed for the algorithm was implemented by bootstrapping the data and running the R package bnlearn (Scutari [2010]¹²) on each bootstrap. The resulting ensemble of networks can be combined to obtain a consensus network where only the most confident edges are kept. Similarly, the estimates of coefficients can be obtained by averaging them over bootstraps.

For this example, the consensus network created with edges with confidence larger than 40% recovers the true structure, and the root mean square error (RMSE) in the coefficient estimates is 0.06 (not shown). This represents a lower bound on the error that we can expect to obtain under perfect reconstruction of the latent space.

When the confounders are unobserved, the reconstruction of the network introduces many false edges and results in a four-times-larger RMSE. Figure 4 shows the reconstructed network, where the red edges are the true edges between V_1 and V_2 and the outcome Z .

We ran algorithms 1 and 2 for 10 iterations using PCA to reconstruct the latent space from the residuals with the assumption that the latent variables were source nodes. We then tracked the latent variable reconstruction as well as the error in the coefficient's estimates. Figure 6 shows the adjusted R^2 between each of the true latent variables and the prediction obtained from the estimated latent space across iterations. The lines and error bands are calculated using LOESS. The estimated latent space is predictive of both the latent variables, and the iterative procedure improves the R^2 with respect to U_1 from 0.57 to 0.61 converging in about 5 iterations.

Figure 7 shows the total error in the coefficients between all variables in the networks and the outcome Z (RMSE) as well as the error in the coefficients of the true drivers of Z , V_1 and V_2 . Both errors converge after the first iteration to an error level of the same magnitude as the error when all variables are observed (dashed lines).

Figure 5 shows the final inferred network at iteration 10. The number of edges arriving to the outcome was reduced considerably with respect to the network prior to inferring latent variables (Figure 4). In addition, the coefficients connecting V_1 and V_2 to the outcome are now closer to their true values. This represents an improvement in ACE estimation.

¹algorithm: hill climbing w/100 restarts, prior: vsp, score: bge

²Both the testbed and the underlying R library will be provided as a supplement. The R library implements the use of PCA, robust PCA, and kernel PCA. Regular PCA was used in this example




Figure 3: True network.




Figure 4: Estimated network when U_1 and U_2 are unobserved.




Figure 5: Estimated network at the last iteration of algorithm 1.

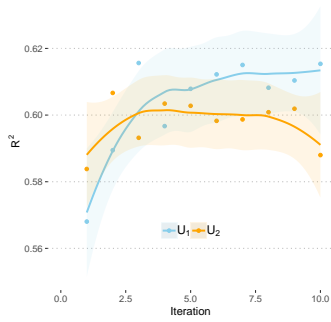


Figure 6: R^2 in the prediction of the latent variable from the selected principal components.

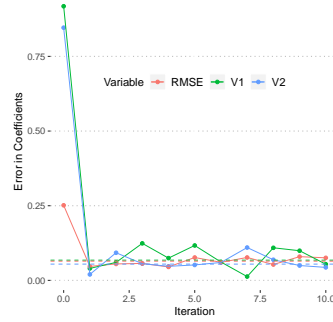


Figure 7: Error in coefficients as a function of the iterations.

7 Conclusions and Future Directions

In this work we present a method for describing the latent variable space that is optimal under linearity, up to rank-linearity. When we cannot provide such guarantees, the method will still identify the terms of the model matrix of the latent space, including any interactions, for models in the GGM family, making it possible to attempt to infer the original (compact) latent space by non-linear modeling and regularization. The method does not place *a priori* constraints on the number of latent variables, and will infer the upper bound on this dimensionality automatically. This method is a generalization of prior work, both in terms of the global treatment of the residual space, and in terms of stronger statements of applicability in cases when probability-scale residuals are applicable.

In the future, we hope to assess the compressibility of the latent space with deep-learning models, using linear PCA to set the ceiling for the cardinality of the latent space. Using autoencoders has been shown to be of high utility when the structure is known (Louizos et al. [2017]). Learning structure over observed data as well as unconstrained latent space via autoencoders would result in a hybrid "deep causal" model that would be useful, for instance, in epidemiological problems. In that space, it is typically highly desirable to retain original features for explainability but also helpful to introduce latent variables to account for unobserved pleiotropic confounding. As noted above, calculation of ACE in epidemiological problems is one specific use case. Applications of such "deep causal networks" should also arise outside of the traditional causal domain - epidemiology. For instance, in computer vision coupled with sensor data, or in the nascent internet-of-things paradigm, retention of some features coupled with the discovery of latent compressed structure may enable inference of cause-and-effect and robust learning.

References

- Animashree Anandkumar, Daniel Hsu, Adel Javanmard, and Sham Kakade. Learning Linear Bayesian Networks with Latent Variables. In Sanjoy Dasgupta and David McAllester, editors, *Proceedings of the 30th International Conference on Machine Learning*, volume 28 of *Proceedings of Machine Learning Research*, pages 249–257, Atlanta, Georgia, USA, June 2013. PMLR. URL <http://proceedings.mlr.press/v28/anandkumar13.html>.
- Elias Bareinboim and Judea Pearl. Causal inference and the data-fusion problem. *Proceedings of the National Academy of Sciences*, 113(27):7345–7352, July 2016. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.1510507113. URL <http://www.pnas.org/lookup/doi/10.1073/pnas.1510507113>.
- Alexander D’Amour. On Multi-Cause Causal Inference with Unobserved Confounding: Counterexamples, Impossibility, and Alternatives. *arXiv:1902.10286 [cs, stat]*, February 2019. URL <http://arxiv.org/abs/1902.10286>. arXiv: 1902.10286.
- Gal Elidan and Nir Friedman. Learning Hidden Variable Networks: The Information Bottleneck Approach. *J. Mach. Learn. Res.*, 6:81–127, December 2005. ISSN 1532-4435. URL <http://dl.acm.org/citation.cfm?id=1046920.1046924>.
- Gal Elidan, Noam Lotner, Nir Friedman, and Daphne Koller. Discovering Hidden Variables: A Structure-Based Approach. In T. K. Leen, T. G. Dietterich, and V. Tresp, editors, *Advances in Neural Information Processing Systems 13*, pages 479–485. MIT Press, 2001. URL <http://papers.nips.cc/paper/1940-discovering-hidden-variables-a-structure-based-approach.pdf>.
- Gal Elidan, Iftach Nachman, and Nir Friedman. “Ideal Parent” Structure Learning for Continuous Variable Bayesian Networks. *Journal of Machine Learning Research*, 8:35, August 2007.
- Nir Friedman. The Bayesian structural EM algorithm. In *Proceedings of the Fourteenth conference on Uncertainty in artificial intelligence*, pages 129–138, 1998. bibtex[organization=Morgan Kaufmann Publishers Inc.].
- Nir Friedman and Daphne Koller. Being Bayesian about Network Structure. *arXiv:1301.3856 [cs, stat]*, January 2013. URL <http://arxiv.org/abs/1301.3856>. arXiv: 1301.3856.
- Nir Friedman and others. Learning belief networks in the presence of missing values and hidden variables. In *ICML*, volume 97, pages 125–133, 1997. bibtex[number=July].
- Matan Gavish and David L. Donoho. The Optimal Hard Threshold for Singular Values is $\sqrt{3}$. *IEEE Transactions on Information Theory*, 60(8):5040–5053, August 2014. ISSN 0018-9448, 1557-9654. doi: 10.1109/TIT.2014.2323359. URL <http://ieeexplore.ieee.org/document/6846297/>.
- Miguel A. Hernán and James M. Robins. Estimating causal effects from epidemiological data. *Journal of Epidemiology and Community Health*, 60(7):578–586, July 2006. ISSN 0143-005X. doi: 10.1136/jech.2004.029496.
- Alexandros Karatzoglou, Alex Smola, Kurt Hornik, and Achim Zeileis. **kernlab** - An S4 Package for Kernel Methods in R. *Journal of Statistical Software*, 11(9), 2004. ISSN 1548-7660. doi: 10.18637/jss.v011.i09. URL <http://www.jstatsoft.org/v11/i09/>.
- Daphne Koller and Nir Friedman. *Probabilistic graphical models: principles and techniques*. Adaptive computation and machine learning. MIT Press, Cambridge, MA, 2009. ISBN 978-0-262-01319-2.
- Alexander Kraskov, Harald Stögbauer, and Peter Grassberger. Estimating mutual information. *Physical Review E*, 69(6), June 2004. ISSN 1539-3755, 1550-2376. doi: 10.1103/PhysRevE.69.066138. URL <https://link.aps.org/doi/10.1103/PhysRevE.69.066138>.
- Christos Louizos, Uri Shalit, Joris Mooij, David Sontag, Richard Zemel, and Max Welling. Causal Effect Inference with Deep Latent-Variable Models. *arXiv:1705.08821 [cs, stat]*, May 2017. URL <http://arxiv.org/abs/1705.08821>. arXiv: 1705.08821.
- David J. C. MacKay. *Information theory, inference, and learning algorithms*. Cambridge University Press, Cambridge, UK ; New York, 2003. ISBN 978-0-521-64298-9.
- Judea Pearl. *Causality: models, reasoning, and inference*. Cambridge University Press, Cambridge, U.K.; New York, 2000. ISBN 978-1-139-64936-0 978-0-511-80316-1. URL <http://dx.doi.org/10.1017/CB09780511803161>. OCLC: 834142635.

- 331 John O. Rawlings, Sastry G. Pantula, and David A. Dickey. *Applied regression analysis: a research tool*.
 332 Springer texts in statistics. Springer, New York, 2nd ed edition, 1998. ISBN 978-0-387-98454-4.
- 333 Marco Scutari. Learning Bayesian Networks with the **bnlearn** R Package. *Journal of Statistical Software*, 35(3),
 334 2010. ISSN 1548-7660. doi: 10.18637/jss.v035.i03. URL <http://www.jstatsoft.org/v35/i03/>.
- 335 Bryan E. Shepherd, Chun Li, and Qi Liu. Probability-scale residuals for continuous, discrete, and censored
 336 data. *The Canadian journal of statistics = Revue canadienne de statistique*, 44(4):463–479, December 2016.
 337 ISSN 0319-5724. doi: 10.1002/cjs.11302. URL [https://www.ncbi.nlm.nih.gov/pmc/articles/](https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5364820/)
 338 [PMC5364820/](https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5364820/).
- 339 Ricardo Silva, Richard Scheines, Clark Glymour, and Peter Spirtes. Learning the Structure of Linear Latent
 340 Variable Models. *J. Mach. Learn. Res.*, 7:191–246, 2006.
- 341 Peter Spirtes, Clark N. Glymour, and Richard Scheines. *Causation, prediction, and search*. Number 81 in
 342 Lecture notes in statistics. Springer-Verlag, New York, 1993. ISBN 978-0-387-97979-3.
- 343 Yixin Wang and David M. Blei. The Blessings of Multiple Causes. *arXiv:1805.06826 [cs, stat]*, May 2018.
 344 URL <http://arxiv.org/abs/1805.06826>. arXiv: 1805.06826.